

Sep 21, 22 23:14	main.py	Page 1/3
<pre>#!/bin/env python3.8 """ Rosemary Cho Based on Example assignment. Author: Chris Curro """ import os import logging import matplotlib import numpy as np import pandas as pd from numpy import pi import matplotlib.pyplot as plt from pdb import set_trace import tensorflow as tf from absl import app from absl import flags from tqdm import trange from sklearn.model_selection import train_test_split from sklearn.preprocessing import MinMaxScaler from keras.regularizers import l2 from tensorflow.keras.models import Sequential from tensorflow.keras.layers import Dense, Flatten import keras import csv from dataclasses import dataclass, field, InitVar script_path = os.path.dirname(os.path.realpath(__file__)) @dataclass class Data: def __post_init__(self): """Dataset source: https://www.kaggle.com/datasets/oddrational/mnist-in-csv?select=mnist_test.csv""" # 60,000 images df_train = pd.read_csv("mnist_train.csv") # 10,000 images df_test = pd.read_csv("mnist_test.csv") X = df_train[df_train.columns[1:]] X_test = df_test[df_test.columns[1:]] y = df_train["label"] y_test = df_test["label"] X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.2, random_state=1) scaler = MinMaxScaler(feature_range=(0, 1)).fit(X_train) X_train = scaler.transform(X_train) X_val = scaler.transform(X_val) X_test = scaler.transform(X_test) X_train = X_train.reshape(X_train.shape[0], 28, 28) X_test = X_test.reshape(X_test.shape[0], 28, 28) X_val = X_val.reshape(X_val.shape[0], 28, 28) X_train = np.reshape(X_train, (-1, 784)) X_test = np.reshape(X_test, (-1, 784)) X_val = np.reshape(X_val, (-1, 784))</pre>		

Sep 21, 22 23:14	main.py	Page 2/3
<pre> self.x_val = X_val self.y_val = y_val self.x_test = X_test self.y_test = y_test self.x = X_train self.y = y_train def get_val_test(self): return self.x_val, self.y_val, self.x_test, self.y_test def get_train(self): X = self.x y = self.y.values return X, y font = { "size": 10, } matplotlib.style.use("classic") matplotlib.rc("font", **font) FLAGS = flags.FLAGS flags.DEFINE_integer("batch_size", 16, "Number of samples in batch") flags.DEFINE_float("learning_rate", 1e-3, "Learning rate / step size for SGD") flags.DEFINE_integer("num_iters", 300, "Number of SGD iterations") def main(a): logging.basicConfig() data = Data() input = tf.keras.Input(shape=(784,), name="digits") x1 = tf.keras.layers.Dense(128, activation="relu")(input) x2 = tf.keras.layers.Dense(128, activation="relu")(x1) x3 = tf.keras.layers.Dropout(0.05)(x2) x4 = tf.keras.layers.Dense(64, activation="relu", kernel_regularizer=l2(0.001))(x3) output = tf.keras.layers.Dense(10, name="predictions")(x4) model = tf.keras.Model(inputs=input, outputs=output) optimizer = tf.optimizers.Adam(learning_rate=FLAGS.learning_rate) lossfunc = tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True) train_metric = keras.metrics.SparseCategoricalAccuracy() val_metric = keras.metrics.SparseCategoricalAccuracy() test_metric = keras.metrics.SparseCategoricalAccuracy() X_train, y_train = data.get_train() X_val, y_val, X_test, y_test = data.get_val_test() """Reference: https://keras.io/guides/writing_a_training_loop_from_scratch/""" train_dataset = tf.data.Dataset.from_tensor_slices((X_train, y_train)) train_dataset = train_dataset.shuffle(buffer_size=1024).batch(FLAGS.batch_size) val_dataset = tf.data.Dataset.from_tensor_slices((X_val, y_val)) val_dataset = val_dataset.batch(FLAGS.batch_size)</pre>		

Sep 21, 22 23:14

main.py

Page 3/3

```

train_log = []
val_log = []
test_log = []

bar = trange(FLAGS.num_iters)

for i in bar:
    # one iteration -> training across entire dataset
    for step, (x_batch, y_batch) in enumerate(train_dataset):
        with tf.GradientTape() as tape:
            logits = model(x_batch, training=True)
            loss = lossfunc(y_batch, logits)

            train_metric.update_state(y_batch, logits)
            train_acc = train_metric.result()
            train_log.append(train_acc.numpy())

            val_logits = model(X_val, training=False)
            val_metric.update_state(y_val, val_logits)
            val_acc = val_metric.result()
            val_metric.reset_states()
            val_log.append(val_acc.numpy())

        grads = tape.gradient(loss, model.trainable_weights)
        optimizer.apply_gradients(zip(grads, model.trainable_weights))
        bar.set_description(f"Loss @ {i} => {loss.numpy():0.6f}")
        bar.refresh()

test_logits = model(X_test, training=False)
test_metric.update_state(y_test, test_logits)
test_acc = test_metric.result()
test_log.append(test_acc.numpy())

plt.title("Model Accuracy")
plt.ylabel("Accuracy")
plt.xlabel("Epoch")
plt.plot(train_log, label="Train", c="blue")
plt.plot(val_log, label="Val", c="purple")
plt.axhline(y=0.955, linestyle="--", c="green", label="0.955")
if test_log[0] >= 0.955:
    plt.plot(len(train_log), test_log[0], marker="*", label="Test")
plt.legend(loc="lower right")
plt.xlim(0, len(train_log) + 200)
plt.ylim(0.3, 1)

plt.tight_layout()
plt.savefig(f"{script_path}/mnist.pdf")

if __name__ == "__main__":
    app.run(main)

# ./tf.sh python mnist/main.py --batch_size 32 --num_iters 1
# ./tf.sh gs -sDEVICE=pdfwrite -dNOPAUSE -dBATCH -dSAFER -sOutputFile=mnist.pdf
mnist/main.py.pdf mnist/mnist.pdf

```

Model Accuracy

